

Implementation of Parallel Quick Sort using MPI

*CSE 633: Parallel Algorithms
Dr. Russ Miller*


*Deepak Ravishankar Ramkumar
50097970*

Recap of Quick Sort

- Given a list of numbers, we want to sort the numbers in increasing or decreasing order.
- On a single processor the unsorted list is in its primary memory, and at the end the same processor would contain the sorted list in its primary memory.
- Quicksort is generally recognized as the fastest sorting algorithm based on comparison of keys, in the average case.
- Quicksort has some natural concurrency.

Sequential quicksort algorithm:

- Select one of the numbers as pivot element.
- Divide the list into two sub lists: a “low list and a “high list”
- The low list and high list recursively repeat the procedure to sort themselves.
- The final sorted result is the concatenation of the sorted low list, the pivot, and the sorted high list



Algorithm for Parallel Quick Sort

- Start off assuming that the number of processors are a power of two.
- At the completion of the algorithm,
 - (1) the list stored in every processor's memory is sorted, and
 - (2) the value of the last element on processor P_i is less than or equal to the value of the first element on P_{i+1} .


Parallel quicksort algorithm

- Randomly choose a pivot from one of the processes and broadcast it to every process.
- Each process divides its unsorted list into two lists: those smaller than (or equal) the pivot, those greater than the pivot
- Each process in the upper half of the process list sends its “low list” to a partner process in the lower half of the process list and receives a “high list” in return.
- The processes divide themselves into two groups and the algorithm is recursive.


- Now, the upper-half processes have only values greater than the pivot, and the lower-half processes have only values smaller than the pivot.
- After $\log P$ recursions, every process has an unsorted list of values completely disjoint from the values held by the other processes.
- The largest value on process i will be smaller than the smallest value held by process $i + 1$.
- Each process now can sort its list using sequential quicksort.




(a)




(b)




(c)



(d)



(e)




(f)

Readings for 1 Million Numbers

	A	B	C	D	E	F	G	H
1	Num of Processors	Reading 1	Reading 2	Reading 3	Reading 4	Reading 5	Average	
2	1	0.486137	0.479304	0.498915	0.459345	0.481175	0.48098	
3	2	0.405326	0.406831	0.407008	0.412961	0.403531	0.40713	
4	4	0.337033	0.357055	0.333783	0.342317	0.333216	0.34068	
5	8	0.291766	0.280861	0.296402	0.281548	0.287413	0.2876	
6	16	0.36568	0.374642	0.362062	0.365493	0.358666	0.36531	
7	32	0.649314	0.298465	0.435157	0.272094	0.278392	0.38668	
8	64	0.621614	0.578988	0.612795	0.633433	0.620268	0.61342	
9	128	1.09468	1.24556	1.1024	1.0973	1.21765	1.15152	
10	256	1.96445	1.87924	1.79648	1.77215	1.71713	1.82589	
11								

Nodes vs time 1 Million Data Items

Time




Readings for 10 Million Numbers

	A	B	C	D	E	F	G	H
1	Num of Processors	Reading 1	Reading 2	Reading 3	Reading 4	Reading 5	Average	
2	1	5.63154	7.55708	5.59229	5.2965	5.61848	5.93918	
3	2	3.63402	3.65614	3.60379	3.7325	3.69719	3.66473	
4	4	3.08463	3.08168	3.19005	3.17345	3.18753	3.14347	
5	8	3.23096	2.84125	2.91948	3.2931	2.86998	3.03095	
6	16	2.79442	4.19635	4.16732	7.66447	4.58758	4.68203	
7	32	4.09503	7.22516	6.9171	8.04961	8.15057	6.88749	
8	64	10.9961	5.66727	15.8492	5.90791	9.70833	9.62576	
9	128	9.76145	10.1217	10.9627	9.4038	12.8083	10.6116	
10	256	15.4053	15.7564	15.4962	15.616	15.5478	15.5643	
11								

Nodes vs time 10 Million Data Items

■ Time




Readings for 50 Million Numbers

G1 : ✕ ✓ f _x Average								
	A	B	C	D	E	F	G	H
1	Num of Processors	Reading 1	Reading 2	Reading 3	Reading 4	Reading 5	Average	
2	1	30.9067	28.4611	28.1991	28.4471	30.9821	29.3992	
3	2	19.6237	19.1829	20.4774	19.1715	19.1127	19.5136	
4	4	23.9516	23.3604	22.8615	22.5307	22.6656	23.074	
5	8	14.1024	13.6956	13.6535	13.7426	13.7267	13.7842	
6	16	12.3328	12.1905	12.1798	12.1843	12.1892	12.2153	
7	32	11.5322	11.1226	11.2357	11.0757	10.9717	11.1876	
8	64	24.3776	24.0729	24.2584	24.5778	24.1861	24.2946	
9	128	39.7297	39.2354	39.0507	39.0437	39.8945	39.3908	
10	256	78.6667	79.8285	78.2345	80.0237	78.7782	79.1063	
11								

Nodes vs time

50 Million Data Items

Time




Readings for 100 Million Numbers

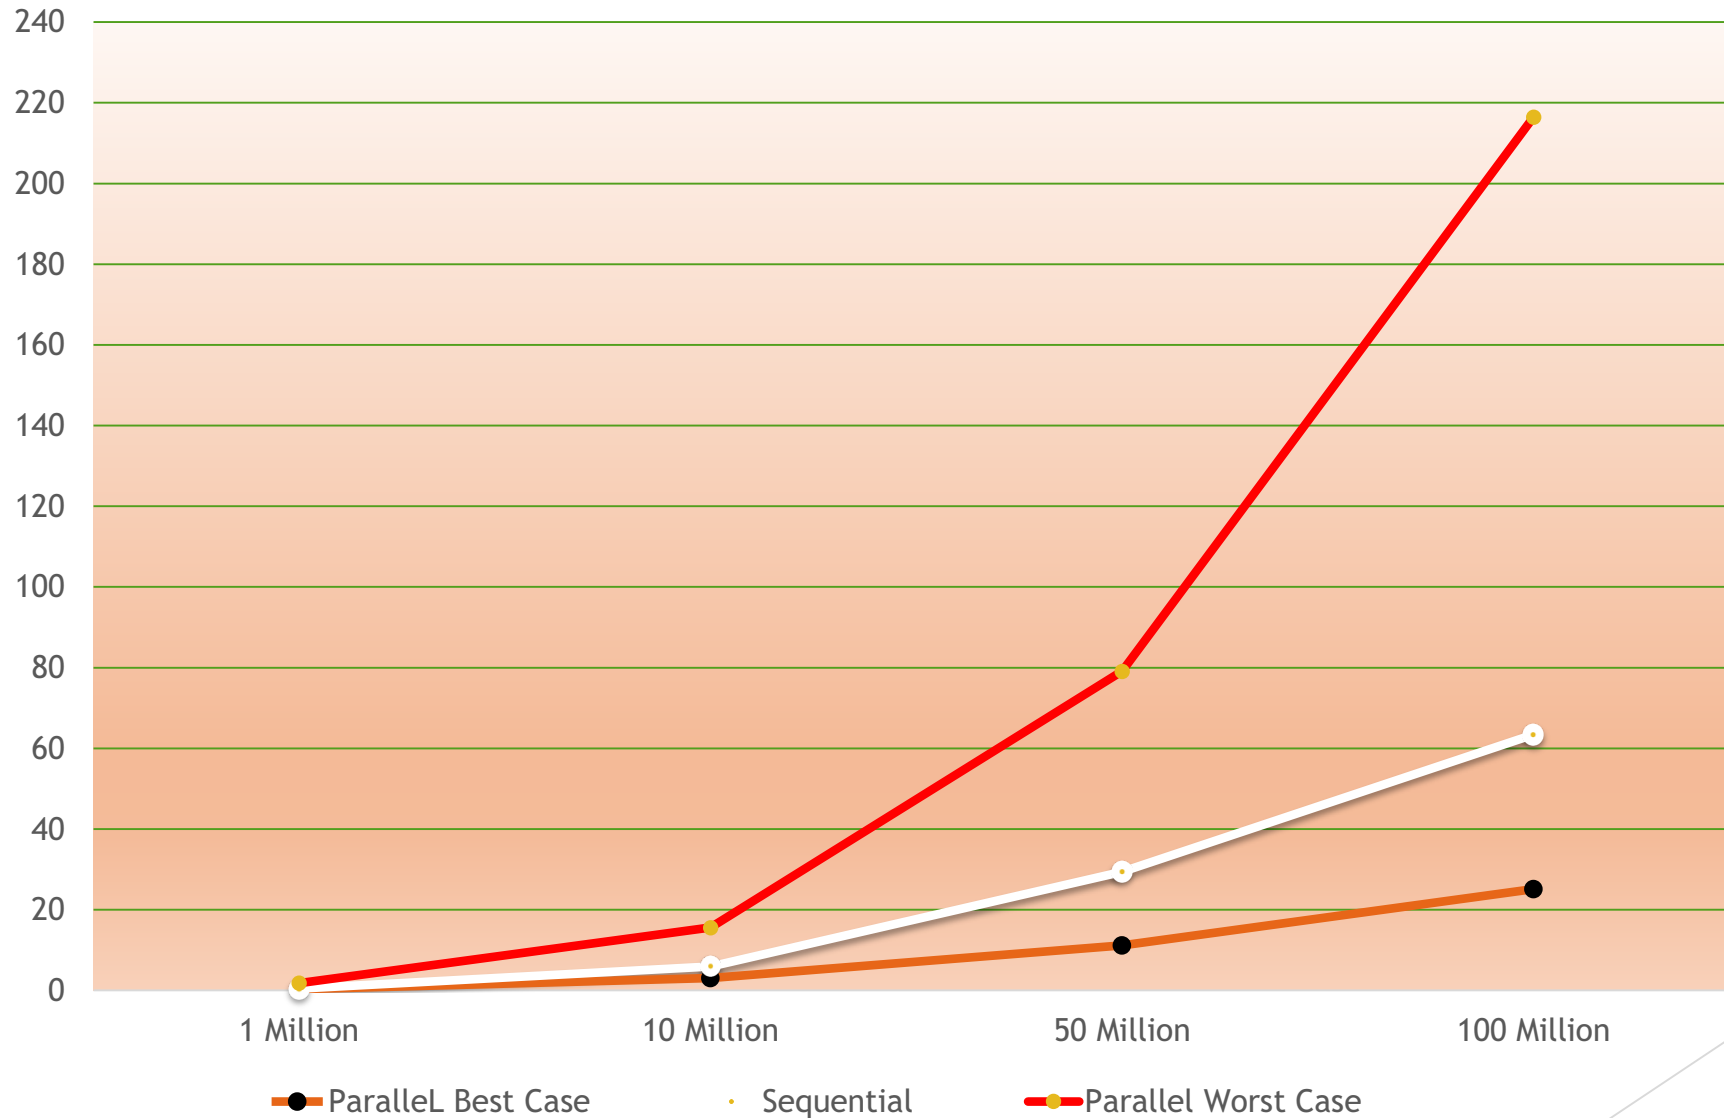
	A	B	C	D	E	F	G	H
1	Num of Processors	Reading 1	Reading 2	Reading 3	Reading 4	Reading 5	Average	
2	1	69.3125	60.4611	60.4283	62.8341	63.9237	63.3919	
3	2	51.3782	50.4934	50.5294	52.3765	50.2241	51.0003	
4	4	44.8693	44.8474	44.7109	44.4998	44.24	44.6335	
5	8	41.5838	38.9771	39.1323	39.1884	39.7165	39.7196	
6	16	31.5832	23.0936	23.1187	24.7625	23.1056	25.1327	
7	32	41.4321	37.7752	36.2429	38.8732	37.6523	38.3951	
8	64	65.2272	66.792	68.235	66.145	67.2456	66.729	
9	128	113.61	126.206	128.456	117.64	121.743	121.531	
10	256	214.938	234.765	254.26	203.567	174.852	216.476	
11								

Nodes vs time 100 Million Data Items

■ Time



Sequential vs Parallel Speed up



Take away

- This parallel quicksort algorithm is likely to do a poor job of load balancing.
- Even if one processor has work to do all the other processes have to wait for it to complete.
- Also faced deadlock problems and had to make sure that the blocking functions in MPI were used correctly.
- In order to achieve better performance its critical to identify the optimal number of processors that would be required for any given computation.

References

- Miller Algorithms Sequential and Parallel A Unified Approach 3rd edition
- Parallel Programming in C with MPI and OpenMP by Michael J. Quinn

The background features abstract, overlapping green geometric shapes, primarily triangles and polygons, in various shades of green, ranging from light lime to dark forest green. These shapes are concentrated on the right side of the slide, with some extending towards the left. The overall effect is a modern, layered design.

► *Thank you.*